

What the Public Wants Regulated: User Preferences on Online Political Advertising in Six Policy Domains

Anonymous for review

October 5, 2026

Abstract

In the last decade, platforms and policymakers have written and rewritten the rules for online political advertising. However, to date, research on these governance questions has focused on what policies platforms and governments adopt, while the policies that citizens actually want remain unexplored. We examine ordinary user preferences for online political ad regulation across six policy domains that have driven real regulatory debates: misinformation, the scope of permissible ads, news-organization exemptions, special protections for voting information, sponsor disclosure, and microtargeting. Our data come from a nationally representative probability sample of U.S. adults fielded through NORC’s AmeriSpeak Panel in November–December 2020 ($N = 1,061$). We find four patterns. First, trust in platforms to handle political advertising is low across the partisan spectrum, with Republicans expressing more distrust than Democrats. Second, support for regulation varies sharply across domains. Bipartisan support is strong for sponsor disclosure, limits on microtargeting, and removing news-organization exemptions. Views are divided on misinformation removal and on protections for voting information. Third, and relatedly, partisan gaps are concentrated in a few topics rather than diffuse, and largest on items involving misinformation. Finally, when respondents evaluate real ads drawn from the 2020 Facebook Ad Library, approval patterns vary systematically with the content and partisan valence of the ads, with some appeals generating broad consensus and others eliciting pronounced partisan disagreement. Our findings show that some regulations have cross-partisan public support. For others, especially those involving misinformation, any platform choice, including inaction, will draw partisan objection.

1 Introduction

Online political advertising has been a central site of regulatory experimentation since the 2016 and 2020 U.S. election cycles, and the pace of change has only increased. Major platforms have repeatedly altered their policies, with changes ranging from complete bans on political advertising to limits on microtargeting and partial disclosure requirements. More recently, platforms have modified or withdrawn political ad services in response to evolving legal requirements (Chee, 2025), while parallel debates have emerged over how to handle AI-generated political content and fact-checking (Klepper, 2023). Across these shifts, the same core questions recur: which categories of political ads should be permitted, what information about them must be disclosed, and which ads and advertisers should fall within the scope of platform rules.

Scholarly attention to these issues has mostly tracked the supply side, focusing on which policies are enacted and how they operate in practice. Prior research has documented how platforms design and enforce political-ad policies (Kreiss and McGregor, 2019, 2018), how transparency archives work in practice (Edelson et al., 2020; Pochat et al., 2022; Edelson et al., 2019), and how regulators and courts have approached the gap between commercial-speech doctrine and digital political advertising (Fowler et al., 2021). What ordinary users *want* regulated — the demand side of technology governance — has received much less attention. The existing demand-side research is generally limited either to a narrow scope (such as examining only attitudes toward microtargeting; McCarthy 2020; Knight Foundation and Gallup 2020) or comparing general attitudes on online advertising to views on more traditional ads on broadcast media (Ridout et al., 2024).

Yet, research on how the public views the core features of possible regulatory frameworks for online advertising is important for several different stakeholders. For the scholarly community, a description of demand-side preferences is a precondition for evaluating whether observed platform behavior aligns with, diverges from, or leads public opinion. For platforms, which act as *de facto* private regulators of political speech, knowing where public opinion is

divided helps anticipate the shape and size of backlash that any given policy choice is likely to generate. For policymakers, rules that lack public support are more fragile and vulnerable to legal challenge or political contestation. And for democratic debate more broadly, platform political-ad policy sits at the intersection of free expression, electoral integrity, and privacy. Normative arguments about how any of those should be weighted are difficult to have without a clearer description of what the relevant public actually wants.

In this article, we describe public preferences for the regulation of political advertising on social media across six concrete policy domains that have structured real regulatory debates: misinformation, the scope of permissible political ads, news-organization exemptions, special protections for voting information, sponsor disclosure, and microtargeting. Each of these domains corresponds to policies platforms adopted during the 2020 US election cycle. Several also correspond to measures the EU has since codified or to focal points of recent U.S. state and federal legislation. Within these domains, we ask four descriptive questions. First, how much do users trust platforms to handle political advertising appropriately, and how does this baseline vary by partisanship? Second, how does support for regulation vary across the six domains: where is the public broadly supportive, broadly opposed, or divided? Third, where is partisan disagreement concentrated, and where is it minimal? Fourth, do users' abstract regulatory preferences track their judgments of real political ads drawn from the 2020 Facebook Ad Library?

Our data come from a high-quality, nationally representative probability sample of U.S. adults fielded through NORC's AmeriSpeak Panel in November–December 2020. We find low trust in platforms across all parties, with Republicans more distrustful than Democrats, a sharp variation in support across the six domains, partisan disagreement that concentrates on misinformation and voting rather than spreading evenly, and real-ad evaluations that largely (but not always) match the abstract attitudes the public expresses.

Taken together, these patterns suggest that the regulatory landscape divides into two kinds of cases. In some domains, the public offers clear cross-partisan backing that can

ground rules claiming broad legitimacy. In others, particularly around misinformation, any choice, whether action or inaction, will draw partisan objection. In such situations, platforms and regulators will need to justify their choices based on considerations beyond public preferences — such as protecting democratic integrity, ensuring election security, and reducing real-world harm.

2 Background

Online political advertising is contested across platform policy teams, legislatures, and courts. Rules are being written, challenged, revised, and reversed in close succession, and the stakes are large: online ads accounted for roughly 18 percent of U.S. political ad spending in the 2020 cycle, nearly \$1.6 billion (Homonoff, 2020), and the share has grown since.

Yet, most of what scholars know about this regulatory churn comes from the supply side — which policies platforms adopt, how they enforce them, and how they have changed. Far less is known about the demand side — what users actually want regulated, and where their preferences match or diverge from the rules being written. Both platforms and legislators claim public backing for the rules they write, but what the public actually wants is poorly mapped. Describing those preferences across the specific policy dimensions on the table is a precondition for evaluating whether any rule, platform-driven or statutory, can claim broad public support.

2.1 Supply-side and demand-side research on ad governance

Most scholarly work on online political-ad regulation studies the rules, the rule-makers, or the effectiveness of platform policies. Qualitative and institutional research documents that Facebook, Google, and Twitter act as active regulators of paid political speech, developing detailed and frequently revised policies about which ads may run, who may run them, and

what information must be disclosed (e.g., Kreiss and McGregor, 2018, 2019).¹ Researchers have also studied the transparency tools platforms have built around these policies: ad archives have enabled aggregate analysis of political advertising for the first time (Edelson et al., 2019), though coverage and enforcement remain uneven (Edelson et al., 2020; Pochat et al., 2022).

Direct evidence on what users actually *want* from political-ad regulation is relatively thin, but it sits within a larger literature on attitudes toward online advertising and platform-managed political content. The literature documents persistent discomfort with targeting, tracking, and personalization, with users consistently wanting greater transparency, more control, and clearer limits on how their data are used (Boerman et al., 2017; McDonald and Cranor, 2010; Aguirre et al., 2015). On political content specifically, Pew and Knight/Gallup surveys document low confidence in tech companies and bipartisan concern about platform influence on elections (Auxier, 2020; Anderson and Auxier, 2020; Knight Foundation and Gallup, 2020), and survey-experimental work shows discomfort with targeted political advertising on privacy grounds (Baum et al., 2019). These studies establish broad unease but do not show how opinions vary across the specific regulatory dimensions platforms and governments are actively contesting beyond limits on microtargeting.

A small recent literature speaks more directly to online political-ad regulation. Ridout et al. (2024) compare U.S. adults’ views of online versus television political advertising and find online ads are viewed somewhat more favorably. Zhu et al. (2025) analyzes a UK survey in which respondents evaluated real 2019 Facebook ads, finding that citizens do not view political ads as inherently unacceptable and that acceptability is shaped by partisanship and content. Zeng et al. (2021a,b) document a structured user vocabulary for “bad” online ads—clickbait, untrustworthy, distasteful—and show that political ads in the 2020 U.S. cycle exhibited many of the same dark patterns, including bait-and-switch poll formats. Edelson et al. (2025) goes further upstream and asks how the public implicitly defines “political” in

¹Note that these same firms also advise campaigns on how to run those ads effectively, which complicates any picture of platforms as neutral arbiters.

the first place, showing that the public’s definition does not match the definitions platforms have adopted. The mismatch matters because ads classified as non-political fall outside every downstream rule.

However, what this combined literature does not yet offer is a direct description of Americans’ preferences across the regulatory dimensions actively in play. Nor does it check whether stated preferences for abstract rules track judgments of concrete ads.

2.2 Six domains, four questions

We examine six regulatory domains where platforms and governments are actively writing rules: misinformation, the scope of permissible ads, news-organization exemptions, voting-related content, sponsor disclosure, and microtargeting. These dimensions together motivate four questions about how user preferences are structured that we articulate here before presenting our data and results. At least one major platform had implemented each of these rules by 2020, which is why they were included. Nevertheless, these and related questions remain the subject of ongoing debate.

We begin with trust, the precondition for accepting platform-driven rules of any kind. Pew and Knight/Gallup surveys document low public confidence in technology companies’ handling of political content and bipartisan concern about platform influence on elections (Auxier, 2020; Anderson and Auxier, 2020; Knight Foundation and Gallup, 2020). Whether that distrust is symmetric across parties is a separate question, and one that matters in its own right. If trust in platforms is itself partisan, support for any rule platforms enforce will be partisan too.

RQ1. *How much do users trust platforms to handle political advertising appropriately, and how does that baseline vary by partisanship?*

The six regulatory domains we examine were active sites of platform and government rule-making when the survey was fielded, and remain so today. Misinformation policy has cycled

through aggressive enforcement and partial reversal, with Meta’s 2025 termination of its U.S. third-party fact-checking program the most recent inflection (Kaplan, 2025). Platforms have also disagreed on the more basic question of whether political ads belong at all. Twitter banned candidate and issue advertising in late 2019 (Conger, 2019), while Facebook and Google maintained broad permissive policies. News-organization exemptions are pervasive across major platforms (Facebook, 2022; X (formerly Twitter), 2022). Voting-related content is treated in two distinct ways. Platforms permit and at times encourage civic-participation messaging, but restrict ads that spread misinformation about voting procedures, eligibility, or outcomes. Sponsor disclosure is a central feature of proposals like the Honest Ads Act (Klobuchar, 2018). Microtargeting has been the focus of its own federal proposal, the Banning Microtargeted Political Ads Act (Zacharia, 2020).

Each domain has been studied largely in isolation, and no prior work has mapped public preferences across them side by side. The shape of that variation, where support is broad, where it is narrow, and where the public is divided, determines which rules can claim broad legitimacy.

RQ2. *How does support for regulation vary across the six domains?*

A separate question concerns the partisan structure of these preferences. Theoretical work on motivated reasoning, affective polarization, and identity-based belief predicts that partisans evaluate political content asymmetrically (Taber and Lodge, 2006; Iyengar and Westwood, 2015), and adjacent empirical work on content moderation and ad evaluation finds clear partisan structure (Kozyreva et al., 2023; Appel et al., 2023; Pradel et al., 2024). Direct evidence on political-ad regulation specifically is thinner, and we treat partisanship as exploratory throughout. The open question is where disagreement concentrates, since concentrated partisan disagreement implies specific battles over specific rules, while diffuse disagreement implies any rule will be contested.

RQ3. *Where is partisan disagreement concentrated, and where is it minimal?*

A final question concerns whether users’ commitments to abstract rules survive contact with specific ads. Endorsing a rule in principle is one thing, but reacting to a real ad that the rule would permit or prohibit is another. We therefore also ask for respondents’ judgments about real Facebook ads to see whether and where the two come apart.

RQ4. *Do users’ stated regulatory preferences track their judgments of concrete ads?*

Together, these four questions structure the empirical work that follows. We turn next to the survey instrument and the four sets of analyses it supports.

3 Data and methods

Our data come from a wave of NORC’s AmeriSpeak Panel fielded in November–December 2020 ($N = 1,061$). AmeriSpeak is a probability sample of U.S. adults recruited from area-probability and address-based sampling frames. Post-stratification weights provided by NORC are applied throughout. Sampling, recruitment, panel composition, and coding decisions are documented in Appendix A. Table 1 reports the weighted composition of the analytic sample.

3.1 Survey instrument

The instrument administered three blocks relevant to the present paper (other items were included for unrelated studies). We summarize each below. Full item wording for all three blocks is reproduced in Appendix B.

Trust. Five items measured trust in mass media, in Facebook content generally, and in Facebook, Twitter, and Google to handle political advertising. All five items used a four-point response (1 = great deal, 2 = fair amount, 3 = not very much, 4 = none at all). For analysis the scale is reversed so that higher values indicate more trust.

Table 1: Weighted sample composition.

Variable	Category	Pct
Party (3-way)	Democrat	47.3
	Independent	15.7
	Republican	37.0
Age	18-29	20.7
	30-44	25.1
	45-59	24.3
	60+	30.0
Gender	Male	48.5
	Female	51.5
Education	Less than HS	9.8
	HS	27.8
	Some college	27.6
	BA	21.3
	Postgrad	13.5
Race/ethnicity	White, non-Hisp.	65.3
	Black, non-Hisp.	12.4
	Other, non-Hisp.	1.7
	Hispanic	17.3
	2+, non-Hisp.	3.3

Notes: Weighted percentages from NORC's AmeriSpeak Panel, fielded November–December 2020 ($N = 1,061$). Cells show the share of weighted respondents in each category.

Regulation preferences. Twenty policy-preference items asked respondents how strongly they agreed with regulatory statements about political advertising on social media. Responses were recorded on a six-point scale (1 = strongly agree, 6 = strongly disagree). Items were written in both pro-regulation and anti-regulation directions. Figures 2–4 display weighted means on a 1–6 scale where 6 = strongly agree and 1 = strongly disagree with the stated item position, so that items people tend to support appear toward the high end of the axis regardless of how they were worded. For domain composites (reported in Appendix C), items are first harmonized so that higher values consistently indicate stronger support for regulation; the 13 pro-regulation-worded items are reverse-coded by subtracting from 7. The 20 items map onto our six policy domains as listed in Table 2. Item presentation order was randomized within respondents.

Real ads. A third block presented respondents with seven ads drawn from the 2020 Facebook Ad Library, shown in randomized order. For each ad, respondents were asked whether Facebook should allow ads of that kind on a four-point scale from definitely allow to definitely not allow. We report per-ad results as the share saying Facebook should allow the ad (the top two response options collapsed). The stimulus set is small and unbalanced across direction $\times$ extremity $\times$ content type. The two Republican-leaning ads are both flagged as misinformation, and there is no Republican-leaning ad in the moderate or non-misinformation cell. We discuss the implications of this design constraint in Section 4.4. Per-ad characteristics, including partisan direction, extremity, and a misinformation flag, are reported in Table 5. Ad images are reproduced in Appendix B.

3.2 Analytic strategy

We report survey-weighted means with 95% confidence intervals computed via the `survey` package in R (Lumley, 2004). Partisan analyses use a three-way classification based on a standard seven-point party identification measure (codes 1–3 as Democrats, 4 as Indepen-

Table 2: Mapping of the 20 policy-preference items to the six domains, with item wording and direction.

Domain	Item	Wording	Direction
Misinformation	ADS4_A	Remove ads with false info	Pro-regulation
Misinformation	ADS4_B	Allow ads even if false	Anti-regulation
Misinformation	ADS4_C	Label false ads with warning	Pro-regulation
Misinformation	ADS4_D	Don't label false ads	Anti-regulation
Misinformation	ADS4_P	Remove candidate misinfo	Pro-regulation
Misinformation	ADS4_Q	Allow all candidate ads	Anti-regulation
Scope of ads	ADS4_E	No political ads at all	Pro-regulation
Scope of ads	ADS4_F	Candidates: no restriction	Anti-regulation
Scope of ads	ADS4_G	Remove controversial-topic ads	Pro-regulation
Scope of ads	ADS4_S	Ban candidate ads, allow issue	Pro-regulation
Scope of ads	ADS4_T	Ban PAC and dark-money ads	Pro-regulation
News orgs	ADS4_I	Exempt news organizations	Anti-regulation
News orgs	ADS4_J	News orgs: same rules	Pro-regulation
Voting	ADS4_H	Allow voting encouragement	Pro-regulation
Voting	ADS4_K	Ban ads discrediting voting	Pro-regulation
Sponsor disclosure	ADS4_L	Disclose funding	Pro-regulation
Sponsor disclosure	ADS4_M	No funding disclosure required	Anti-regulation
Sponsor disclosure	ADS4_R	Public ad database	Pro-regulation
Microtargeting	ADS4_N	Ban microtargeting	Pro-regulation
Microtargeting	ADS4_O	Allow audience selection	Anti-regulation

Notes: The 20 policy-preference items, mapped to the six policy domains. “Direction” indicates whether the item is reverse-coded for analysis (Pro-regulation: agreement expresses support for regulation, reverse-coded so that high values run pro-regulation; Anti-regulation: agreement expresses opposition to regulation, kept as-is). Full item wording appears in Appendix B.

dents, 5–7 as Republicans). Differences between Democrats and Republicans are tested with survey-weighted two-group t -tests. Within each family of comparisons (the 20 policy-preference items, the 7 real ads, the 5 trust items, and the 6 domain composites used in Section 4.2), p -values are adjusted using the Holm method to account for multiple testing. We treat the four families as separate because they map to distinct research questions, not because we assume independence between families. The resulting per-family correction is less conservative than a global correction across all comparisons would be, and we make this choice deliberately to retain power for the per-domain comparisons. Significance markers reported in tables and figures throughout the paper use Holm-adjusted p -values.

For domain-level summaries, we construct composites by averaging each respondent’s harmonized scores across items within the domain. We use these composites in Section 4.2 to characterize support and the partisan pattern at the domain level. Item-level results follow in tables and figures.

4 Results

We present four findings, one for each research question stated in Section 2.2.

4.1 Trust in platforms is low and partisan-shaped (RQ1)

Trust in mass media and in social-media platforms to handle political advertising is low across all parties. Figure 1 reports weighted mean trust on the four-point reversed scale (1 = none at all, 4 = great deal) for each of the five items by party. Appendix Table 7 reports the same numbers with 95% confidence intervals and Democrat–Republican differences.

On the mass-media item, the overall weighted mean is 2.53. Democrats average 2.91, Independents 2.39, and Republicans 2.11. The Democrat–Republican difference is 0.79 points (95% CI [0.67, 0.91], Holm-adjusted $p < 0.05$). The four platform items run lower still. Trust in Facebook content in general is 1.84 overall, with no significant Democrat–Republican




Figure 1: Weighted mean trust by party. Reversed scale: 1 = none at all, 4 = great deal. Error bars are 95% confidence intervals. An asterisk on a measure label indicates a Holm-adjusted Democrat–Republican difference at $p < 0.05$. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

gap. Trust in Facebook, Twitter, and Google to handle political advertising specifically averages 1.73, 1.77, and 2.15. Across these three platform-handling items, Democrats are less distrustful than Republicans. The Democrat–Republican gaps are 0.40 for Facebook, 0.49 for Twitter, and 0.64 for Google, all Holm-significant ($p < 0.05$).

Two patterns are worth noting. First, no group on average expresses even “a fair amount” of trust in any of the four platform items. The highest single platform mean is Democratic trust in Google (2.41), still below the midpoint of the scale. Distrust is the modal stance everywhere. Second, the partisan structure of distrust runs in one direction. Republicans are more distrustful of mass media and of platform handling of political advertising than Democrats are. The pattern is most pronounced for Google (Democrat–Republican gap of 0.64) and least pronounced for general Facebook content, where the parties are statistically indistinguishable.

4.2 Support for regulation varies sharply across the six domains (RQ2)

Support for regulation varies sharply across the six policy domains. Figure 2 reports overall weighted means for each of the 20 policy-preference items, grouped by domain. Figure 3 adds the partisan breakdown for the same items. Item-level means by party appear in Table 3, and domain-level composite means in Appendix Table 8.

On the 1–6 scale where 6 = strongly agree, the six domains divide into two groups. Three domains show strong consensus and small partisan differences. Three show more divided opinion and pronounced partisan gaps.

Bipartisan domains. Sponsor disclosure is the most strongly supported domain in the battery. The composite mean is 4.99 for Democrats and 4.81 for Republicans, with no Holm-significant difference. The component items run consistently in the pro-regulation direction. Respondents agree with requiring funding disclosure (mean 5.02), disagree with exempting from disclosure requirements (mean 2.02), and agree with maintaining a public ad database (mean 4.60). All three items show negligible partisan separation.

News-organization rules show similar bipartisan opposition to current platform exemptions. Composite means are 4.80 for Democrats and 4.80 for Republicans (gap 0.00, $p > 0.9$). The two component items move together: respondents oppose exempting news organizations from political-ad rules (mean 2.45) and support holding news-organization ads to the same rules as other ads (mean 5.03).

Microtargeting limits draw bipartisan support at a more moderate level. The domain composite is 4.15 for Democrats and 4.03 for Republicans (gap 0.11, not significant). Both items fall in the moderate range of the 1–6 scale with negligible partisan separation: respondents lean toward supporting a microtargeting ban (mean 3.87) and away from permitting unrestricted audience selection (mean 2.62).

Across the three bipartisan domains, the strength of consensus varies. Sponsor disclosure




Figure 2: Overall weighted mean for each of the 20 policy-preference items, on the 1–6 scale where 6 = strongly agree and 1 = strongly disagree with the stated item position. Items grouped by domain and ordered within domain by mean response. Error bars are 95% confidence intervals. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).




Figure 3: Partisan breakdown of the 20 policy-preference items, on the 1–6 scale where 6 = strongly agree and 1 = strongly disagree with the stated item position, grouped by domain. An asterisk on an item label indicates a Holm-adjusted Democrat–Republican difference at $p < 0.05$. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

and news-organization items show the strongest consensus, with means toward the ends of the 1–6 scale. Microtargeting items sit in the moderate range, though the cross-partisan agreement holds.

Divided domains. Misinformation shows large within-domain partisan separation. The composite mean is 4.73 for Democrats and 3.99 for Republicans, with a gap of 0.74 (95% CI [0.59, 0.89], Holm-adjusted $p < 0.05$). Five of the six misinformation items show large partisan gaps that survive Holm correction, with the largest at ‘Companies should not try to label political ads as false, even if they contain misleading information’ (gap 1.17). The exception is ‘Political ads with false content should be allowed, but labeled with a warning,’ where Democrats and Republicans converge toward the disagree end of the scale ($D = 2.68$, $R = 2.74$). This item was designed to mimic the policy put in place by Facebook in advance of the 2020 election, but the wording allows for two-sided disagreement. Democrats agree with a policy of misinformation removal more strongly than Republicans (‘Political ads should be removed if they contain false information’: $D = 5.30$, $R = 4.38$), and disagree more with the proposition that companies should not label false ads (‘Companies should not try to label political ads as false, even if they contain misleading information’: $D = 2.27$, $R = 3.44$).

Voting protections is the second divided domain. The composite is 4.90 for Democrats and 4.02 for Republicans (gap 0.87, Holm-adjusted $p < 0.05$). The split runs entirely through ‘Ads that discredit voting by mail should not be permitted on social media,’ where the partisan gap reaches 1.47 points on the 1–6 scale ($D = 4.65$, $R = 3.18$, Holm-adjusted $p < 0.05$). This is the largest item-level partisan gap in the entire 20-item battery. The other voting item, ‘Ads that encourage voting should always be permitted,’ is bipartisan ($D = 5.15$, $R = 4.89$). Voting-related misinformation, on the evidence here, divides along the same partisan lines as misinformation more broadly.

Scope of permissible ads is divided to a smaller degree, largely because both parties exhibit opinions closer to the middle of the spectrum. The composite is 3.68 for Democrats

Table 3: Item-level weighted means by party, all 20 policy-preference items.

Domain	Item	Wording	Overall	Democrat	Independent	Republican	D-R sig
Misinformation	A	Remove ads with false info	4.82	5.30	4.42	4.38	*
Misinformation	B	Allow ads even if false	2.17	1.79	2.30	2.59	*
Misinformation	C	Label false ads with warning	2.71	2.68	2.70	2.74	
Misinformation	D	Don't label false ads	2.80	2.27	2.84	3.44	*
Misinformation	P	Remove candidate misinfo	4.80	5.22	4.47	4.40	*
Misinformation	Q	Allow all candidate ads	2.12	1.75	2.19	2.54	*
Scope of ads	E	No political ads at all	3.26	3.24	3.30	3.27	
Scope of ads	F	Candidates: no restriction	2.54	2.02	2.62	3.16	*
Scope of ads	G	Remove controversial-topic ads	2.94	3.07	3.14	2.68	
Scope of ads	S	Ban candidate ads, allow issue	2.49	2.51	2.61	2.41	
Scope of ads	T	Ban PAC and dark-money ads	4.38	4.61	4.26	4.13	*
News orgs	I	Exempt news organizations	2.45	2.50	2.37	2.42	
News orgs	J	News orgs: same rules	5.03	5.10	4.84	5.03	
Voting	H	Allow voting encouragement	4.97	5.15	4.59	4.89	
Voting	K	Ban ads discrediting voting	4.00	4.65	3.98	3.18	*
Sponsor disclosure	L	Disclose funding	5.02	5.19	4.66	4.98	
Sponsor disclosure	M	No funding disclosure required	2.02	1.90	2.10	2.12	
Sponsor disclosure	R	Public ad database	4.60	4.70	4.37	4.58	
Microtargeting	N	Ban microtargeting	3.87	3.89	4.12	3.73	
Microtargeting	O	Allow audience selection	2.62	2.61	2.52	2.67	

Notes: Weighted mean responses to each policy-preference item on the 1–6 scale where 6 = strongly agree and 1 = strongly disagree with the stated item position. An asterisk in the final column indicates a Holm-adjusted Democrat–Republican difference at $p < 0.05$ within the 20-item policy-preference family. See Table 4 for items ranked by absolute Democrat–Republican gap and Appendix B for full item wording. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

and 3.27 for Republicans (gap 0.41, Holm-adjusted $p < 0.05$). The gap is driven by ‘Political candidates should always be allowed to post ads without restriction’ (D = 2.02, R = 3.16, gap 1.14). Democrats are much more comfortable than Republicans imposing some kind of restriction on candidate ads. Given the other items in the battery related to misinformation, it may be that these attitudes enter into responses given the phrase “without restriction.” The other scope items show smaller or non-significant differences. Only ‘Ads from certain sources like political action committees (PACs) and dark money organizations should be banned’ (gap 0.47) is significant after Holm correction.

Together, the six-domain map is fairly clear. The three domains where current regulatory debate concerns transparency, targeting practices, or specialized content categories show high cross-partisan support. The three domains where current debate involves judgments about truth or about the legitimate scope of political speech show pronounced partisan separation.

4.3 Partisan disagreement concentrates on misinformation and voting (RQ3)

Where partisan disagreement appears in the battery, it concentrates rather than spreading evenly across domains. Figure 4 reports the Democrat and Republican mean for each of the 20 policy-preference items, ranked within domain by absolute gap size. Holm-significant gaps are marked on item labels and indicated by darker connecting segments. Table 4 ranks items by absolute gap size.

Eight of the 20 items show partisan gaps significant at the Holm-adjusted 0.05 level. Five are misinformation items. One is the voting-related misinformation item (‘Ads that discredit voting by mail should not be permitted on social media’). The remaining two are scope items: the unrestricted candidate ad item just discussed and ‘Ads from certain sources like political action committees (PACs) and dark money organizations should be banned’. The 12 non-significant items spread across the other domains: all three sponsor-disclosure items, both microtargeting items, both news-organization items, the remaining scope items,




Figure 4: Democrat and Republican weighted means for each of the 20 policy-preference items, on the 1–6 scale where 6 = strongly agree and 1 = strongly disagree with the stated item position. Connecting segments drawn in dark grey indicate Holm-adjusted Democrat–Republican differences at $p < 0.05$; light grey indicates non-significant differences. Items grouped by domain. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

Table 4: Items ranked by absolute Democrat–Republican gap.

Domain	Item	Wording	D	R	D - R	Sig
Voting	K	Ban ads discrediting voting	4.65	3.18	1.47 [1.19, 1.76]	*
Misinformation	D	Don't label false ads	2.27	3.44	1.17 [0.91, 1.43]	*
Scope of ads	F	Candidates: no restriction	2.02	3.16	1.14 [0.90, 1.38]	*
Misinformation	A	Remove ads with false info	5.30	4.38	0.92 [0.68, 1.16]	*
Misinformation	P	Remove candidate misinfo	5.22	4.40	0.81 [0.59, 1.04]	*
Misinformation	B	Allow ads even if false	1.79	2.59	0.80 [0.56, 1.03]	*
Misinformation	Q	Allow all candidate ads	1.75	2.54	0.79 [0.55, 1.03]	*
Scope of ads	T	Ban PAC and dark-money ads	4.61	4.13	0.47 [0.23, 0.72]	*
Scope of ads	G	Remove controversial-topic ads	3.07	2.68	0.39 [0.12, 0.65]	
Voting	H	Allow voting encouragement	5.15	4.89	0.26 [0.06, 0.45]	
Sponsor disclosure	M	No funding disclosure required	1.90	2.12	0.22 [0.00, 0.44]	
Sponsor disclosure	L	Disclose funding	5.19	4.98	0.21 [0.03, 0.39]	
Microtargeting	N	Ban microtargeting	3.89	3.73	0.16 [0.00, 0.43]	
Sponsor disclosure	R	Public ad database	4.70	4.58	0.12 [0.00, 0.32]	
Scope of ads	S	Ban candidate ads, allow issue	2.51	2.41	0.09 [0.00, 0.32]	
News orgs	I	Exempt news organizations	2.50	2.42	0.08 [0.00, 0.31]	
News orgs	J	News orgs: same rules	5.10	5.03	0.07 [0.00, 0.26]	
Microtargeting	O	Allow audience selection	2.61	2.67	0.06 [0.00, 0.31]	
Misinformation	C	Label false ads with warning	2.68	2.74	0.06 [0.00, 0.34]	
Scope of ads	E	No political ads at all	3.24	3.27	0.03 [0.00, 0.32]	

Notes: The 20 policy-preference items, ranked by absolute Democrat–Republican gap. Means are weighted means on the 1–6 scale, where 6 = strongly agree and 1 = strongly disagree with the stated item position. The $|D - R|$ column reports the absolute partisan gap; brackets show the 95% confidence interval. An asterisk in the final column indicates a Holm-adjusted Democrat–Republican difference at $p < 0.05$ within the 20-item policy-preference family. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

and the bipartisan voting item (‘Ads that encourage voting should always be permitted’).

The items where partisan disagreement is largest cluster around two themes. The misinformation items and the voting-related item involve respondent judgments about truth and about which truth claims should be regulated. The item calling for fully unrestricted candidate ads (item F), whose gap puts it third in the ranking, can reasonably be interpreted as falling in the same category.

4.4 Real-ad evaluations track abstract preferences (RQ4)

The seven real ads drawn from the 2020 Facebook Ad Library track the abstract battery in some categories cleanly and in others not at all. Figure 5 reports the share of respondents saying Facebook should allow each ad, by party. Table 5 reports the same numbers with Democrat–Republican differences and Holm-adjusted significance markers, alongside per-ad characteristics.

Three descriptive patterns appear. First, bipartisan civic content draws bipartisan approval. The neutral get-out-the-vote ad is approved by 92% of respondents overall, with Democrats at 92% and Republicans at 93%. The Democrat-leaning moderate GOTV ad draws 80% allow, with no significant Democrat–Republican difference. Where the content is bipartisan or civic, the parties agree.

Second, partisans approve own-side ads more than other-side ads. Republicans approve the two Republican-leaning ads (the stolen-votes claim and the antifa-burning-cities ad) at 67% and 65%, while Democratic approval of the same two ads is 22% and 29% (gaps of 45 and 36 points, both Holm-significant). Democrats approve the moderate Democrat-leaning ad at 83% and the convict-Trump ad at 58%, while Republican approval is 78% and 44% (the convict-Trump gap is Holm-significant). On these items, in-group preference is a driving feature, although both parties prefer the more extreme political ad less.

Third, Democrats appear more willing than Republicans to reject own-side ads that fail standards of truthfulness or appropriateness. The two Republican-leaning ads draw




Figure 5: Weighted share saying Facebook should allow each of the seven real ads, by party. An asterisk on an ad label indicates a Holm-adjusted Democrat–Republican difference at $p < 0.05$. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

Table 5: Per-ad approval rates and Democrat–Republican gaps.

Ad	Direction	Extremity	Misinfo	V-misinfo	Content	All	Dem	Ind	Rep	D-R	Sig
ADS5	D-leaning	extreme	No	No	Police brutality (lib. media)	40%	35%	34%	48%	-0.12 [-0.21, -0.04]	*
ADS6	D-leaning	extreme	No	No	Cops as pigs (org)	34%	35%	35%	32%	0.03 [-0.05, 0.11]	
ADS7	R-leaning	–	Yes	Yes	Stolen-votes claim (cons. media)	41%	22%	33%	67%	-0.45 [-0.53, -0.37]	*
ADS8	Bipartisan	–	No	No	GOTV neutral	92%	92%	89%	93%	-0.01 [-0.06, 0.03]	
ADS9	D-leaning	moderate	No	No	GOTV (liberal)	80%	83%	78%	78%	0.05 [-0.02, 0.12]	
ADS11	D-leaning	extreme	No	No	Convict Trump	51%	58%	44%	44%	0.14 [0.05, 0.23]	*
ADS12	R-leaning	extreme	Yes	No	Antifa cities (org)	45%	29%	45%	65%	-0.36 [-0.44, -0.27]	*

Notes: Approval is the weighted share of respondents saying Facebook should allow the ad (top two of the four-point allow scale). Columns “All”, “Dem”, “Ind”, and “Rep” show the share for the overall sample, Democrats, Independents, and Republicans. “V-misinfo” is the voting-misinformation flag. Brackets in the “D-R” column show the 95% confidence interval on the Democrat–Republican difference in approval shares. An asterisk in the final column indicates a Holm-adjusted Democrat–Republican difference at $p < 0.05$ within the seven-ad family. “Direction”, “Extremity”, “Misinfo”, and “V-misinfo” are author codings of each ad based on the questionnaire’s source and content notes; full per-ad descriptions appear in Appendix B. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

uniformly high Republican approval (67% and 65%), even though both are flagged as misinformation in our coding. The four Democrat-leaning ads, none flagged as misinformation but varying in extremity and tone, draw Democratic approval that ranges widely: 83% for the moderate GOTV ad, 58% for the convict-Trump ad, and 35% each for the police-brutality ad and the anti-police satire. On the two most provocative Democrat-leaning ads, Democratic approval is at or below Republican approval (police-brutality ad: $D = 35\%$, $R = 48\%$, Holm-significant; anti-police satire: $D = 35\%$, $R = 32\%$, not significant). Restricting to the four extreme partisan ads in the stimulus set, Democratic approval averages 43%, below the Republican rate of 66% on the two extreme Republican-leaning ads.

We report this third pattern descriptively. The design does not let us interpret it cleanly since both Republican-leaning ads in the stimulus set are misinformation, and none of the four Democrat-leaning ads are. We cannot tell from this design whether Democrats hold own-side ads to higher truthfulness or appropriateness standards in general, whether Republicans would reject own-side ads if any in the set were misinformation, or whether Democrats would approve own-side misinformation if any were present. A balanced design across direction $\times$ extremity $\times$ misinformation would be needed to separate these.

However, one within-domain comparison is possible. The stolen-votes claim ad directly invokes the abstract policy item ‘Ads that discredit voting by mail should not be permitted on social media,’ so we can compare the abstract preference and the concrete judgment within the same domain. Figure 6 reports the share taking the protective stance on each side: support for the abstract ban on the left, disapproval of the actual ad on the right. Both are partisan-structured in the same direction. The weighted Pearson correlation between support for banning ads that discredit voting and disapproval of the stolen-votes claim ad is $r = 0.38$ overall ($n = 950$), with within-party correlations of $r = 0.22$ for Democrats, $r = 0.30$ for Independents, and $r = 0.28$ for Republicans (Table 6).




Figure 6: Convergent evidence on voting misinformation. Left panel: weighted share supporting a ban on ads that discredit voting (top three of the harmonized scale). Right panel: weighted share saying Facebook should not allow the stolen-votes claim ad. Error bars are 95% confidence intervals. Both panels show the same protective-stance direction. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

Table 6: Abstract and concrete responses to voting misinformation, by party. Left panel: weighted share supporting a ban on ads that discredit voting (top three of the harmonized scale). Right panel: weighted share saying Facebook should not allow the stolen-votes claim ad. Error bars are 95% confidence intervals. Both panels show the same protective-stance direction. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

Group	Weighted r	n
Overall	0.382	950
Democrat	0.220	485
Independent	0.300	121
Republican	0.277	344

Notes: Weighted Pearson correlation between agreement with the item banning ads that discredit voting (harmonized so that higher values indicate stronger support for the ban) and disapproval of the stolen-votes claim ad (recoded so that higher values indicate that the respondent says Facebook should not allow the ad). Both measures are aligned in the protective-stance direction. Sample sizes (n) reflect respondents with non-missing values on both items and a non-missing party identification. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

5 Discussion

The findings divide the regulatory landscape into two kinds of cases. In some domains, the public offers cross-partisan backing for specific rules, and platforms and policymakers can ground their decisions in something close to public consensus. In others, the parties’ preferences diverge sharply, and any rule a platform or government adopts will draw partisan objection from one side or the other. Knowing which kind of case a regulatory decision falls into is consequential. It tells regulators whether they are moving with public opinion, against it, or into partisan contested ground.

Three of the six policy domains are bipartisan in our data. Sponsor disclosure draws strong support across the partisan spectrum. Microtargeting limits draw support at a more moderate level, but with the same structure of agreement. Subjecting news-organization ads to the same rules as other political ads is supported by both parties, and existing exemptions (Facebook, 2022; X (formerly Twitter), 2022) appear to lack public legitimacy. Each of these is a domain where current platform or legal rules either match public preferences or could be tightened with cross-partisan support. For these domains, public opinion is a resource available to platforms and policymakers who choose to mobilize it.

The other three domains are partisan-divided. Misinformation regulation produces large within-domain partisan separation, mirroring trends in moderation-preferences research more broadly (Kozyreva et al., 2023; Pradel et al., 2024). Voting protections show the same structure. A single survey item about banning ads that discredit voting produces the largest item-level partisan gap in our entire battery, larger than any item in the misinformation domain. The scope-of-permissible-ads domain is divided to a smaller degree, primarily through the question of whether candidates’ ads should run without restriction. In each of these three domains, no choice a platform makes will satisfy both Democrats and Republicans. Removing misinformation will be coded as censorship by some Republican elites and accepted by Democratic elites. Leaving it in place will be coded as inaction by Democrats and as legitimate restraint by Republicans. The same dynamic applies to voting-related content.

5.1 Recommendations for platforms

Platforms set the rules for political advertising on their services, and the legitimacy of those rules depends in part on whether they reflect what users actually want. Gaining a better understanding of those user preferences was the purpose of this research, and three concrete changes follow from the bipartisan preferences identified above.

Remove news organization exemptions from political ad rules. Respondents of both parties oppose exempting news organizations and support holding news-organization ads to the same rules as other political ads, with no significant partisanship gap on either item. Existing exemptions on Meta’s platforms (Facebook, 2022; X (formerly Twitter), 2022) and parallel carve-outs elsewhere run directly counter to bipartisan public preference.

Maintain and strengthen sponsor disclosure and public ad archives. In our battery, transparency, both in the form of sponsor disclosure and publicly available political ad archives, enjoys very strong bipartisan support. The recommendation here is partly affirmative; platforms that currently provide ad libraries and require sponsor disclosure should keep doing so. The recommendation is also a brake on recent reversals. Platforms that have narrowed access to ad archives, deprecated transparency tools, or reduced disclosure requirements are moving against a clear, cross-partisan public preference.

Restrict microtargeting of political ads. Bipartisan support for restricting microtargeting is real but more moderate than for disclosure or news-org parity. This indicates directional support rather than the strong consensus seen for disclosure, but it is still a domain where platforms restricting microtargeting can claim cross-partisan U.S. public support.

5.2 Limitations

These findings come with key limitations that should also be considered. First, the data come from a single survey wave fielded in November–December 2020. The questions in the instrument are not specific to the 2020 election. They ask about general regulatory principles, not about specific incidents. The political environment at the time was distinctive, falling

in the immediate post-election period, at the height of claims about voter fraud in the 2020 election, and during a moment of intense platform action on misinformation. We argue above that the underlying questions have continued through the 2025 reversals on EU and U.S. content moderation, but a direct test would require additional waves.

Second, the seven real ads in the stimulus set are unbalanced across direction $\times$ extremity $\times$ content. The set contains no Republican-leaning moderate ad and no Republican-leaning ad outside the misinformation category. This shapes which comparisons the design supports. We can characterize per-ad approval and within-domain convergence (especially for the voting case), but we cannot cleanly separate partisan in-group acceptance from partisan acceptance of misinformation.

Third, in regard to our recommendations in Section 5.1, the survey measures the direction of preferences in the abstract, not the tradeoffs users would accept in practice. Respondents support disclosure, news-organization parity, and microtargeting limits when asked about each in isolation. We do not know how those preferences would shift if respondents were asked to weigh them against advertiser compliance burdens, campaigns' ability to reach relevant voters, or press-freedom concerns associated with subjecting news organizations to the same rules as other political advertisers.

5.3 Conclusion

Regulating political speech under conditions of high partisanship is genuinely hard. On the divided domains, especially misinformation, no platform or government choice will satisfy both sides, and legitimacy will have to come from grounds other than user preferences, such as democratic integrity, electoral security, or harm reduction. The trust findings sharpen the difficulty. Republican distrust of platforms and Republican opposition to misinformation rules likely reflect the same underlying skepticism about platforms as arbiters of truth, which means the harder problem on these rules is not designing them but securing legitimacy for the institution that would enforce them.

Even so, paths remain. Three of the six domains we examine attract cross-partisan public support, and rules in those domains can claim a legitimacy that the divided rules cannot. Identifying which domains fall on which side of that line requires description of what the public actually wants, domain by domain. Description of public preferences is not by itself a basis for policy. It is, however, a precondition for evaluating which policies can plausibly claim public support, which cannot, and where the obstacles to policy legitimacy actually sit. This paper provides that description, mapping public preferences across six domains currently under active regulatory consideration, characterizing where partisan disagreement concentrates, and showing how abstract attitudes line up with judgments of real ads.

References

- Aguirre, E., Mahr, D., Grewal, D., de Ruyter, K., and Wetzels, M. (2015). Unraveling the personalization paradox: The effect of information collection and trust-building strategies on online advertisement effectiveness. *Journal of Retailing*, 91(1):34–49.
- Anderson, M. and Auxier, B. (2020). 55% of U.S. social media users say they are ‘worn out’ by political posts and discussions. <https://www.pewresearch.org/short-reads/2020/08/19/55-of-u-s-social-media-users-say-they-are-worn-out-by-political-posts-and-discussions/>.
- Appel, R. E., Pan, J., and Roberts, M. E. (2023). Partisan conflict over content moderation Is more than disagreement about facts. *Science Advances*, 9(44):eadg6799.
- Auxier, B. (2020). 54% of Americans say social media companies shouldn’t allow any political ads. <https://www.pewresearch.org/short-reads/2020/09/24/54-of-americans-say-social-media-companies-shouldnt-allow-any-political-ads/>.
- Baum, K., Meissner, S., Abramova, O., and Krasnova, H. (2019). Do they really care about targeted political ads? Investigation of user privacy concerns and preferences. In *Proceedings of the 27th European Conference on Information Systems (ECIS)*, Stockholm–Uppsala, Sweden.
- Boerman, S. C., Kruikemeier, S., and Zuiderveen Borgesius, F. J. (2017). Online behavioral advertising: A literature review and research agenda. *Journal of Advertising*, 46(3):363–376.
- Chee, F. Y. (2025). Meta to halt political advertising in EU from October, blames EU rules. <https://www.reuters.com/sustainability/boards-policy-regulation/meta-halt-political-advertising-eu-october-blames-eu-rules-2025-07-25/>.

Conger, K. (2019). Twitter will ban all political ads, C.E.O. Jack Dorsey says. <https://www.nytimes.com/2019/10/30/technology/twitter-political-ads-ban.html>.

Edelson, L., Lauinger, T., and McCoy, D. (2020). A security analysis of the Facebook ad library. In *2020 IEEE Symposium on Security and Privacy*, SP '20, pages 661–678.

Edelson, L., Lockett, D., Guillard, C., Lauinger, T., Li, Z., Montgomery, J. M., and McCoy, D. (2025). What drives perceptions of the political in online advertising?: The source, content, and political orientation. *Journal of Experimental Political Science*, 13(1):36–50.

Edelson, L., Sakhuja, S., Dey, R., and McCoy, D. (2019). An Analysis of United States Online Political Advertising Transparency. arXiv:1902.04385 [cs.SI].

Facebook (2022). Ad authorization exemptions and how they work. <https://www.facebook.com/business/help/387111852028957>.

Fowler, E. F., Franz, M. M., and Ridout, T. N. (2021). *Political Advertising in the United States*. Routledge, New York, 2 edition.

Homonoff, H. (2020). 2020 political ad spending exploded: Did it work? <https://www.forbes.com/sites/howardhomonoff/2020/12/08/2020-political-ad-spending-exploded-did-it-work/?sh=366fa13b3ce0>.

Iyengar, S. and Westwood, S. J. (2015). Fear and loathing across party lines: New evidence on group polarization. *American Journal of Political Science*, 59(3):690–707.

Kaplan, J. (2025). More speech and fewer mistakes. <https://about.fb.com/news/2025/01/meta-more-speech-fewer-mistakes/>.

Klepper, D. (2023). To help 2024 voters, Meta says it will begin labeling political ads that use AI-generated imagery. <https://www.pbs.org/newshour/politics/to-help-2024-voters-meta-says-it-will-begin-labeling-political-ads-that-use-ai-generated>

- Klobuchar, A. (2018). S.1989 – 115th Congress (2017-2018): Honest Ads Act. <https://www.congress.gov/bill/115th-congress/senate-bill/1989>.
- Knight Foundation and Gallup (2020). How Americans view tech and internet companies’ impact on the 2020 election. <https://legacy.knightfoundation.org/articles/how-americans-view-tech-and-internet-companies-impact-on-the-2020-election/>.
- Kozyreva, A., Herzog, S. M., Lewandowsky, S., Hertwig, R., Lorenz-Spreen, P., Leiser, M., and Reifler, J. (2023). Resolving content moderation dilemmas between free speech and harmful misinformation. *Proceedings of the National Academy of Sciences*, 120(7):e2210666120.
- Kreiss, D. and McGregor, S. C. (2018). Technology firms shape political communication: The work of Microsoft, Facebook, Twitter, and Google with campaigns during the 2016 U.S. presidential cycle. *Political Communication*, 35(2):155–177.
- Kreiss, D. and McGregor, S. C. (2019). The “arbiters of what our voters see”: Facebook and Google’s struggle with policy, process, and enforcement around political advertising. *Political Communication*, 36(4):499–522.
- Lumley, T. (2004). Analysis of complex survey samples. *Journal of Statistical Software*, 9(8):1–19.
- McCarthy, J. (2020). In U.S., most oppose micro-targeting in online political ads. <https://news.gallup.com/opinion/gallup/286490/oppose-micro-targeting-online-political-ads.aspx>.
- McDonald, A. M. and Cranor, L. F. (2010). Beliefs and behaviors: Internet users’ understanding of behavioral advertising. In *38th Research Conference on Communication, Information and Internet Policy (TPRC)*.

- Pochat, V. L., Edelson, L., Goethem, T. V., Joosen, W., McCoy, D., and Lauinger, T. (2022). An audit of facebook’s political ad policy enforcement. In *31st USENIX Security Symposium (USENIX Security 22)*, pages 607–624, Boston, MA. USENIX Association.
- Pradel, F., Zilinsky, J., Kosmidis, S., and Theocharis, Y. (2024). Toxic speech and limited demand for content moderation on social media. *American Political Science Review*, 118(4):1895–1912.
- Ridout, T. N., Franz, M. M., Fowler, E. F., and Kang, T. (2024). Public perceptions of television and online political advertising. *Research & Politics*, 11(2).
- Taber, C. S. and Lodge, M. (2006). Motivated skepticism in the evaluation of political beliefs. *American Journal of Political Science*, 50(3):755–769.
- X (formerly Twitter) (2022). How to get exempted as a news publisher from the political content policy. <https://business.x.com/en/help/ads-policies/ads-content-policies/political-content/news-exemption>.
- Zacharia, J. (2020). Congress should pass Rep. Eshoo’s bill restricting microtargeting of political ads. <https://www.sfchronicle.com/opinion/article/Congress-should-pass-Rep-Eshoo-s-bill-15371138.php>.
- Zeng, E., Kohno, T., and Roesner, F. (2021a). What Makes a Bad Ad? User Perceptions of Problematic Online Advertising. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, pages 1–24, Yokohama Japan. ACM.
- Zeng, E., Wei, M., Gregersen, T., Kohno, T., and Roesner, F. (2021b). Polls, clickbait, and commemorative \$2 bills: problematic political advertising on news and media websites around the 2020 U.S. elections. In *Proceedings of the 21st ACM Internet Measurement Conference*, pages 507–525, Virtual Event. ACM.

Zhu, J., Dommett, K., and Stafford, T. (2025). What makes online political ads unacceptable? Interrogating public attitudes to inform regulatory responses. *Humanities and Social Sciences Communications*, 12(1):806.

A Sample and survey details

A.1 The AmeriSpeak Panel

NORC’s AmeriSpeak Panel is a probability-based panel of U.S. households recruited using a multi-stage sampling design that begins with NORC’s National Sample Frame, an area-probability sample of U.S. households. Initial panel recruitment is conducted using a sequential mail, telephone, and in-person protocol, with web access provided to households that would otherwise be excluded by reliance on online recruitment alone. Once enrolled, panel members complete questionnaires online or by phone. The panel is periodically replenished to maintain representativeness of U.S. adults age 18 and older. Post-stratification weighting is applied to align completed survey samples with population benchmarks. Panel-level recruitment and disposition rates are documented in NORC’s published AmeriSpeak methodology.

A.2 This survey wave

The wave used in this paper was fielded by NORC in November–December 2020. The analytic sample is $N = 1,061$ U.S. adults. Post-stratification weights provided by NORC and applied to all weighted analyses in this paper align the sample on age, sex, race/ethnicity, education, and Census region with U.S. adult population benchmarks. The questionnaire administered alongside the items reported here included experimental modules from other research teams that are not analyzed in this paper. The wave was fielded November 16–December 2, 2020. The weighted household recruitment rate (AAPOR RR III for panel recruitment) is 20.4%; the panel retention rate is 80.7%; the survey completion rate is 28.7%; and the weighted cumulative response rate is 4.72%. These figures, along with full disposition codes, are documented in the NORC Card and AmeriSpeak Project Report included in the replication archive.

A.3 Survey instruments

The three blocks of items used in this paper are the trust battery (ADS1, ADS2, ADS3_A through ADS3_C), the 20-item regulation-preferences battery, and the seven real-ad evaluation items (ADS5–ADS9, ADS11, and ADS12). Full item wording for all three blocks is reproduced in Appendix B. Item presentation order was randomized within respondents for each block. Short item labels for the regulation-preferences battery, with domain assignments and direction, appear in Table 2. The seven real-ad images themselves are reproduced in the replication archive.

A.4 Coding decisions

For domain composites, all policy-preference items are harmonized so that higher values indicate stronger support for regulation. The 13 items worded in the pro-regulation direction (A, C, E, G, H, J, K, L, N, P, R, S, T) are reverse-coded by subtracting from 7; the seven anti-regulation items (B, D, F, I, M, O, Q) are kept as-is. Figures 2–4 display means on a 1–6 scale where 6 = strongly agree and 1 = strongly disagree with the stated item position: reversed items retain their harmonized (7 minus raw) values, and the seven non-reversed items are also transformed by subtracting from 7, so that all 20 items run in the same direction for display. The four-point trust scale is reversed so that 4 indicates the highest level of trust and 1 the lowest. Real-ad responses are coded as the share giving one of the top two responses on the four-point allow scale (“definitely allow” or “probably allow”). Domain composites used in §4.2 are constructed by averaging each respondent’s harmonized item scores within the domain, with item-level missingness handled by computing the mean over available items. Respondents missing all items in a domain do not contribute to that domain’s composite.

B Survey item wording

This appendix reproduces the full text of each survey item used in the analyses. The instrument was administered in CAWI (web) format for online respondents and CATI (telephone) format for phone respondents; response options are identical across modes. Within each block, item presentation order was randomized across respondents. Item ADS10 was not part of the fielded instrument; the seven real-ad evaluation items are numbered ADS5 through ADS9 and ADS11 through ADS12.

B.1 Trust block (ADS1–ADS3)

ADS1 (mass media). “In general, how much trust and confidence do you have in the mass media—such as newspapers, TV, and radio—when it comes to reporting the news fully, accurately, and fairly?”

ADS2 (Facebook content). “In general, how much trust and confidence do you have in the information you see on Facebook when it comes to reporting the news fully, accurately, and fairly?”

ADS3 (platform handling of political advertising). Prompt: “Recently, technology companies like Facebook, Twitter, and Google have announced new rules about how political advertisements are handled on their platforms. How much trust and confidence do you have in the following companies to appropriately handle political advertising?” Three items were presented as a grid:

- ADS3_A: Facebook
- ADS3_B: Twitter
- ADS3_C: Google

Response options for all five trust items. A great deal / A fair amount / Not very much / None at all. Coded 1–4 in the raw data; reversed for analysis so that 4 indicates the highest level of trust.

B.2 Regulation-preferences block

Prompt: “How much do you agree with the following statements about how social media companies should handle political advertisements?” Twenty items were presented as a randomized grid.

- ADS4_A: Political ads should be removed if they contain false information.
- ADS4_B: Political ads should be allowed, even if they contain misleading information.
- ADS4_C: Political ads with false content should be allowed, but labeled with a warning.
- ADS4_D: Companies should not try to label political ads as false, even if they contain misleading information.
- ADS4_E: No political ads should be allowed at all.
- ADS4_F: Political candidates should always be allowed to post ads without restriction.
- ADS4_G: All ads on controversial topics should be removed, regardless of source.
- ADS4_H: Ads that encourage voting should always be permitted.
- ADS4_I: Ads from news organizations should be exempted from regulations of political advertising.
- ADS4_J: Political ads from news organizations should be subject to the same regulations as ads from other sources.
- ADS4_K: Ads that discredit voting by mail should not be permitted on social media.

- ADS4_L: Political ads should include information about who paid for them.
- ADS4_M: Political advertisers should not be required to disclose their funding sources.
- ADS4_N: Microtargeting political ads to specific audiences based on things like age, gender, and race should not be permitted.
- ADS4_O: Political advertisers should be allowed to select who sees their ads.
- ADS4_P: Ads that contain misinformation about a candidate should be removed.
- ADS4_Q: All ads about candidates should be allowed, even if they contain false content.
- ADS4_R: Companies should provide a publicly accessible database of all political ads.
- ADS4_S: Ads about candidates or upcoming elections should be banned, but ads about social issues (e.g., guns or abortion) should still be allowed.
- ADS4_T: Ads from certain sources like political action committees (PACs) and dark money organizations should be banned.

Response options for all 20 items. Strongly agree / Agree / Agree only a little / Disagree only a little / Disagree / Strongly disagree. Coded 1–6 in the raw data. Items A, C, E, G, H, J, K, L, N, P, R, S, and T are reverse-coded for analysis so that all 20 items run on a common scale where higher values indicate stronger support for regulation.

B.3 Real-ads block (ADS5–ADS12)


Block prompt: “Below we will show you a series of political advertisements that have appeared on Facebook and ask if you think these kinds of ads should be allowed.”

Each ad was presented as an image followed by the question: “Do you think that Facebook should allow ads like this?” The seven ads were presented in randomized order. Per-ad

content summaries (drawn from the questionnaire’s source and content notes) are listed below; the ad images themselves are reproduced in Figure 7.

- ADS5: Source: liberal media. Content: police brutality.
- ADS6: Source: organization. Content: police-themed satire.
- ADS7: Source: conservative-leaning media. Content: claim that Democrats are stealing votes (coded as voting misinformation in this paper).
- ADS8: Source: neutral. Content: get-out-the-vote ad.
- ADS9: Source: liberal. Content: get-out-the-vote ad.
- ADS11: Source: unknown. Content: urging representatives to convict President Trump.
- ADS12: Source: organization. Content: claim that antifa is burning U.S. cities (coded as misinformation in this paper).

Response options for each ad. Definitely allow / Probably allow / Probably not allow / Definitely not allow. Coded 1–4 in the raw data. Approval reported in §4.4 is the share of respondents giving one of the top two responses.




ADS5




ADS6



ADS7



ADS8



ADS9



ADS11



ADS12

Figure 7: The seven Facebook ads presented in the real-ads block. Each ad was shown to respondents in randomized order with the prompt “Do you think that Facebook should allow ads like this?” on a four-point scale from definitely allow to definitely not allow. Source: ads drawn from the 2020 Facebook Ad Library by the original survey team.

C Supplementary tables

This appendix collects two tables that support the analyses in §4.1 and §4.2 but are not central to the main-text narrative.

Table 7: Weighted mean trust by party, with Democrat–Republican difference (Holm-adjusted within the trust family).

Measure	Democrat	Independent	Republican	Overall	D - R	Sig
Mass media	2.91	2.39	2.11	2.53	0.79 [0.67, 0.91]	*
Facebook (general)	1.89	1.79	1.80	1.84	0.09 [-0.03, 0.22]	
Facebook (pol. ads)	1.91	1.74	1.50	1.73	0.40 [0.28, 0.52]	*
Twitter (pol. ads)	2.01	1.70	1.52	1.77	0.49 [0.36, 0.62]	*
Google (pol. ads)	2.41	2.29	1.77	2.15	0.64 [0.50, 0.78]	*

Notes: Weighted mean trust on the four-point reversed scale (1 = none at all, 4 = great deal). Brackets in the “D - R” column show the 95% confidence interval on the Democrat–Republican difference. An asterisk in the final column indicates a Holm-adjusted Democrat–Republican difference at $p < 0.05$ within the five-item trust battery. Trust items: ADS1 (mass media), ADS2 (Facebook content), ADS3_A–C (Facebook, Twitter, and Google handling of political advertising). Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).

Table 8: Domain-level composite means by party, with Democrat–Republican difference (Holm-adjusted within the six-domain family).

Domain	Democrat	Independent	Republican	D - R	Sig
Misinformation	4.73	4.20	3.99	0.74 [0.59, 0.89]	*
Scope of ads	3.68	3.54	3.27	0.41 [0.25, 0.56]	*
News orgs	4.80	4.74	4.80	-0.00 [-0.18, 0.17]	
Voting	4.90	4.29	4.02	0.87 [0.69, 1.06]	*
Sponsor disclosure	4.99	4.64	4.81	0.18 [0.02, 0.34]	
Microtargeting	4.15	4.29	4.03	0.11 [-0.09, 0.32]	

Notes: Domain composites are the weighted mean of each respondent’s harmonized item scores within the domain on the 1–6 scale (higher values indicate stronger support for regulation). Brackets in the “D - R” column show the 95% confidence interval on the Democrat–Republican difference. An asterisk in the final column indicates a Holm-adjusted Democrat–Republican difference at $p < 0.05$ within the six-domain family. See Appendix B for the items composing each domain. Source: NORC AmeriSpeak Panel, November–December 2020 ($N = 1,061$).